

```
import numpy as np
from .layers import *
from .layer_utils import *

class FullyConnectedNet(object):
    """
    A fully-connected neural network with an arbitrary number of hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be

    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

    where batch normalization and dropout are optional, and the {...} block is
    repeated L - 1 times.

    Similar to the TwoLayerNet above, learnable parameters are stored in the
    self.params dictionary and will be learned using the Solver class.
    """

    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
                  dropout=0, use_batchnorm=False, reg=0.0,
                  weight_scale=1e-2, dtype=np.float32, seed=None):
        """
        Initialize a new FullyConnectedNet.

        Inputs:
        - hidden_dims: A list of integers giving the size of each hidden layer.
        - input_dim: An integer giving the size of the input.
        - num_classes: An integer giving the number of classes to classify.
        - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
          the network should not use dropout at all.
        - use_batchnorm: Whether or not the network should use batch normalization.
        - reg: Scalar giving L2 regularization strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - dtype: A numpy datatype object; all computations will be performed using
          this datatype. float32 is faster but less accurate, so you should use
          float64 for numeric gradient checking.
        - seed: If not None, then pass this random seed to the dropout layers. This
          will make the dropout layers deterministic so we can gradient check the
          model.
        """
        self.use_batchnorm = use_batchnorm
        self.use_dropout = dropout > 0
        self.reg = reg
        self.num_layers = 1 + len(hidden_dims)
        self.dtype = dtype
        self.params = {}

        # ===== #
        # YOUR CODE HERE:
        #   Initialize all parameters of the network in the self.params dictionary.
        #   The weights and biases of layer 1 are W1 and b1; and in general the
        #   weights and biases of layer i are Wi and bi. The
        #   biases are initialized to zero and the weights are initialized
        #   so that each parameter has mean 0 and standard deviation weight_scale.
        #
        #   BATCHNORM: Initialize the gammas of each layer to 1 and the beta
        #   parameters to zero. The gamma and beta parameters for layer 1 should
        #   be self.params['gamma1'] and self.params['beta1']. For layer 2, they
        #   should be gamma2 and beta2, etc. Only use batchnorm if self.use_batchnorm
        #   is true and DO NOT do batch normalize the output scores.
        # ===== #

        # Concat dims for full NN
        dims = [input_dim] + hidden_dims + [num_classes]
        for layer in range(self.num_layers):
            self.params['W' + str(layer + 1)] = np.random.normal(0, weight_scale, (dims[layer], dims[layer + 1]))
            self.params['b' + str(layer + 1)] = np.zeros(dims[layer + 1])

            if self.use_batchnorm and (layer != (self.num_layers-1)):
                self.params['gamma' + str(layer + 1)] = np.ones(dims[layer + 1])
                self.params['beta' + str(layer + 1)] = np.zeros(dims[layer + 1])

        # ===== #
        # END YOUR CODE HERE
        # ===== #

        # When using dropout we need to pass a dropout_param dictionary to each
        # dropout layer so that the layer knows the dropout probability and the mode
        # (train / test). You can pass the same dropout_param to each dropout layer.
        self.dropout_param = {}
        if self.use_dropout:
            self.dropout_param = {'mode': 'train', 'p': dropout}
            if seed is not None:
                self.dropout_param['seed'] = seed

        # With batch normalization we need to keep track of running means and
        # variances, so we need to pass a special bn_param object to each batch
        # normalization layer. You should pass self.bn_params[0] to the forward pass
        # of the first batch normalization layer, self.bn_params[1] to the forward
        # pass of the second batch normalization layer, etc.
        self.bn_params = []
        if self.use_batchnorm:
            self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

        # Cast all parameters to the correct datatype
        for k, v in self.params.items():
            self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since they
    # behave differently during training and testing.
    if self.dropout_param is not None:
        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param[mode] = mode

    scores = None

    # ===== #
    # YOUR CODE HERE:
    #   Implement the forward pass of the FC net and store the output
    #   scores as the variable "scores".
    #
    #   BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
    #   between the affine_forward and relu_forward layers. You may
    #   also write an affine_batchnorm_relu() function in layer_utils.py.
    #
    #   DROPOUT: If dropout is non-zero, insert a dropout layer after
    #   every ReLU layer.
    # ===== #

    a = {}
    norm = {}
    h = {}
    drop = {}
    drop[0] = [X]

    for layer in range(self.num_layers):
        #Affine
        a[layer + 1] = affine_forward(drop[layer][0], self.params['W' + str(layer + 1)], self.params['b' + str(layer + 1)])

        if layer < (self.num_layers-1):
            # BatchNorm
            if self.use_batchnorm: norm[layer + 1] = batchnorm_forward(a[layer + 1][0], self.params['gamma' + str(layer + 1)],
                self.params['beta' + str(layer + 1)], self.bn_params[layer])
            else: norm[layer + 1] = a[layer + 1]
            # ReLU
            h[layer + 1] = relu_forward(norm[layer + 1][0])
            # Dropout
            if self.use_dropout: drop[layer + 1] = dropout_forward(h[layer + 1][0], self.dropout_param)
            else: drop[layer + 1] = h[layer + 1]

    scores = a[self.num_layers][0]

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # If test mode return early
    if mode == 'test':
        return scores

    loss, grads = 0.0, {}
    # ===== #
    # YOUR CODE HERE:
    #   Implement the backwards pass of the FC net and store the gradients
    #   in the grads dict, so that grads[k] is the gradient of self.params[k]
    #   Be sure your L2 regularization includes a 0.5 factor.
    #
    #   BATCHNORM: Incorporate the backward pass of the batchnorm.
    #
    #   DROPOUT: Incorporate the backward pass of dropout.
    # ===== #

    loss, dout = softmax_loss(scores, y)
    Ws = [self.params['W' + str(i + 1)] for i in range(self.num_layers)]

    loss += 0.5 * self.reg * sum([np.linalg.norm(weight, 'fro')**2 for weight in Ws])
    das = {}
    dhs = {}
    ddrops = {}
    dnorms = {}
    dgammas = {}
    dbetas = {}
    dws = {}
    dbs = {}
    das[self.num_layers] = dout

    for layer in reversed(range(self.num_layers)):
        ddrops[layer], dws[layer + 1], dbs[layer + 1] = affine_backward(das[layer + 1], a[layer + 1][1])
        if layer != 0:
            if self.use_dropout: dhs[layer] = dropout_backward(ddrops[layer], drop[layer][1])
            else: dhs[layer] = ddrops[layer]
            dnorms[layer] = relu_backward(dhs[layer], h[layer][1])
            if self.use_batchnorm: das[layer], dgammas[layer], dbetas[layer] = batchnorm_backward(dnorms[layer], norm[layer][1])
            else: das[layer] = dnorms[layer]

    for layer in range(self.num_layers):
        grads['W' + str(layer + 1)] = dws[layer + 1] + self.reg * self.params['W' + str(layer + 1)]
        grads['b' + str(layer + 1)] = dbs[layer + 1].T
        if layer != (self.num_layers-1) and self.use_batchnorm:
            grads['gamma' + str(layer + 1)] = dgammas[layer + 1]
            grads['beta' + str(layer + 1)] = dbetas[layer + 1].T

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grads
```