

```
import numpy as np

from nndl.layers import *
from nndl.conv_layers import *
from utils.fast_layers import *
from nndl.layer_utils import *
from nndl.conv_layer_utils import *

import pdb

class ThreeLayerConvNet(object):
    """
    A three-layer convolutional network with the following architecture:

    conv - relu - 2x2 max pool - affine - relu - affine - softmax

    The network operates on minibatches of data that have shape (N, C, H, W)
    consisting of N images, each with height H and width W and with C input
    channels.
    """

    def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
                  hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
                  dtype=np.float32, use_batchnorm=False):
        """
        Initialize a new network.

        Inputs:
        - input_dim: Tuple (C, H, W) giving size of input data
        - num_filters: Number of filters to use in the convolutional layer
        - filter_size: Size of filters to use in the convolutional layer
        - hidden_dim: Number of units to use in the fully-connected hidden layer
        - num_classes: Number of scores to produce from the final affine layer.
        - weight_scale: Scalar giving standard deviation for random initialization
          of weights.
        - reg: Scalar giving L2 regularization strength
        - dtype: numpy datatype to use for computation.
        """
        self.use_batchnorm = use_batchnorm
        self.params = {}
        self.reg = reg
        self.dtype = dtype

        # ===== #
        # YOUR CODE HERE:
        #   Initialize the weights and biases of a three layer CNN. To initialize:
        #   - the biases should be initialized to zeros.
        #   - the weights should be initialized to a matrix with entries
        #     drawn from a Gaussian distribution with zero mean and
        #     standard deviation given by weight_scale.
        # ===== #
        C, H, W = input_dim
        pad = (filter_size - 1) / 2

        self.params['W1'] = np.random.normal(0, weight_scale, [num_filters,C,filter_size, filter_size])
        self.params['b1'] = np.zeros(num_filters)
        conv_out_h = int((H + 2 * pad - filter_size) + 1)
        conv_out_w = int((W + 2 * pad - filter_size) + 1)
        pool_out_h = int((conv_out_h - 2) / 2 + 1)
        pool_out_w = int((conv_out_w - 2) / 2 + 1)
        self.params['W2'] = np.random.normal(0, weight_scale, [pool_out_h * pool_out_w * num_filters, hidden_dim])
        self.params['b2'] = np.zeros(hidden_dim)
        self.params['W3'] = np.random.normal(0, weight_scale, [hidden_dim, num_classes])
        self.params['b3'] = np.zeros(num_classes)

        # ===== #
        # END YOUR CODE HERE
        # ===== #

        for k, v in self.params.items():
            self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Evaluate loss and gradient for the three-layer convolutional network.

    Input / output: Same API as TwoLayerNet in fc_net.py.
    """
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    W3, b3 = self.params['W3'], self.params['b3']

    # pass conv_param to the forward pass for the convolutional layer
    filter_size = W1.shape[2]
    conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}

    # pass pool_param to the forward pass for the max-pooling layer
    pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

    scores = None

    # ===== #
    # YOUR CODE HERE:
    #   Implement the forward pass of the three layer CNN. Store the output
    #   scores as the variable "scores".
    # ===== #

    if not self.use_batchnorm:
        l1_out, l1_cache = conv_relu_pool_forward(X, W1, b1,conv_param, pool_param)
        l2_out, l2_cache = affine_relu_forward(l1_out, W2, b2)
        scores, l3_cache = affine_forward(l2_out, W3, b3)
    # else:
    #     conv_out, conv_cache = conv_forward_fast(x, w, b, conv_param)
    #     out, cache = spatial_batchnorm_forward(conv_out, gamma, beta, bn_param)
    #     s, relu_cache = relu_forward(a)
    #     out, pool_cache = max_pool_forward_fast(s, pool_param)
    #     out, cache = affine_forward(x, w, b)
    #     out, cache = batchnorm_forward(x, gamma, beta, bn_param)
    #     s, relu_cache = relu_forward(a)
    #     out, cache = affine_forward(x, w, b)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    if y is None:
        return scores

    loss, grads = 0, {}
    # ===== #
    # YOUR CODE HERE:
    #   Implement the backward pass of the three layer CNN. Store the grads
    #   in the grads dictionary, exactly as before (i.e., the gradient of
    #   self.params[k] will be grads[k]). Store the loss as "loss", and
    #   don't forget to add regularization on ALL weight matrices.
    # ===== #

    loss, dout = softmax_loss(scores, y)
    loss += 0.5 * self.reg * sum([np.sum(np.square(weight)) for weight in [W1,W2,W3]])

    dx3, dw3, db3 = affine_backward(dout, l3_cache)
    dx2, dw2, db2 = affine_relu_backward(dx3, l2_cache)
    _, dw1, db1 = conv_relu_pool_backward(dx2, l1_cache)
    grads['b1'], grads['b2'], grads['b3'] = db1, db2, db3
    grads['W1'] = dw1 + self.reg * W1
    grads['W2'] = dw2 + self.reg * W2
    grads['W3'] = dw3 + self.reg * W3

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grads
```