



ECE C147/C247

Introduction to PyTorch

Pan Lu, UCLA CS



Introduction to PyTorch

- What is PyTorch
- PyTorch basics
 - Tensors
 - Neural network layer
- PyTorch CNN example
 - Data preparation
 - Model construction
 - Model training
 - Model evaluation
- PyTorch advances
 - Useful tricks
 - Useful packages and tools



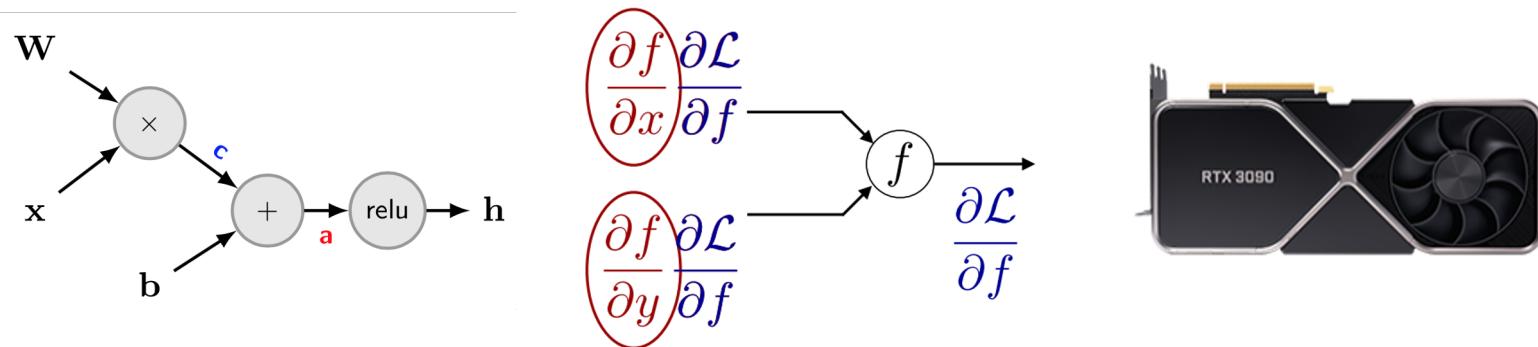
The Colab link for this lecture:

<https://colab.research.google.com/drive/146XJEdXOguMQ41CbU-6TqS2gEuc061ql>



The basic idea of deep learning libraries

- In deep learning libraries, you write code that builds a **computational graph**, akin to what we did in the backpropagation lectures.
- Deep learning libraries **implement backpropagation** for you, so that after the computational graph is built, you don't have to implement the backward pass!!
- Deep learning libraries also provide a straightforward API to do all training on **GPUs**.





A zoo of frameworks!

Caffe
Caffe (UCB)

→  **Caffe2**
Caffe2 (Facebook)

 **torch**
Torch (NYU / Facebook)

→  **PyTorch**
PyTorch (Facebook)

theano
Theano (U Montreal)

→  **TensorFlow**
TensorFlow (Google)

We'll focus on these

 **PaddlePaddle**
PaddlePaddle (Baidu)

 **Chainer**
Chainer

 **mxnet**
MXNet (Amazon)

 **CNTK**
CNTK (Microsoft)

 **JAX** (Google)

And others...



PyTorch

Pytorch is a deep learning framework, i.e... set of functions and libraries which allow you to do higher-order programming designed for Python programming language based on Torch.

Pros:

- Based on Python: pythonic and shares commands with numpy
- Easy to get started: easy to install, simple syntax
- Community: active community, very good documents, including tutorial: <https://pytorch.org/tutorials/>.
- Dynamic computational graphs: the network behavior can be changed programmatically at runtime, easier model optimization
- Data parallelism: work among multiple CPU or GPU cores

Cons:

- Limited monitoring and visualization interfaces
- Not as extensive as TensorFlow



PyTorch

Install PyTorch

- <https://pytorch.org/>
- Mac:
`conda install pytorch torchvision torchaudio -c pytorch`
- Windows:
`conda install pytorch torchvision torchaudio cpuonly -c pytorch`

PyTorch Build	Stable (1.10.2)	Preview (Nightly)	LTS (1.8.2)
Your OS	Linux	Mac	Windows
Package	Conda	Pip	LibTorch Source
Language	Python		C++ / Java
Compute Platform	CUDA 10.2	CUDA 11.3	ROCM 4.2 (beta) CPU
Run this Command:	<code>conda install pytorch torchvision torchaudio -c pytorch</code>		



PyTorch

PyTorch tutorial

- <https://pytorch.org/tutorials/>
- PyTorch provide a very good online tutorials and documents

The screenshot shows the PyTorch tutorial website at the 'Learn the Basics' page. The top navigation bar includes links for Get Started, Ecosystem, Mobile, Blog, Tutorials (which is highlighted), Docs, Resources, and GitHub. A search bar labeled 'Search Tutorials' is also present. The main content area features a large red box highlighting the 'Learn the Basics' section, which contains links to Quickstart, Tensors, Datasets & DataLoaders, Transforms, Build the Neural Network, Automatic Differentiation with `torch.autograd`, Optimizing Model Parameters, and Save and Load the Model. Below this is a heading 'LEARN THE BASICS' with a subtext 'You can play with example codes to learn'. It lists authors: Suraj Subramanian, Seth Juarez, Cassie Breviu, Dmitry Soshnikov, Ari Bornstein. It describes the tutorial's purpose of introducing ML workflows in PyTorch. It mentions the use of the FashionMNIST dataset to train a neural network. A note states that the tutorial assumes basic Python and Deep Learning familiarity.

1.10.1+cu102

PyTorch Recipes [+]

Introduction to PyTorch [-]

Learn the Basics

- Quickstart
- Tensors
- Datasets & DataLoaders
- Transforms
- Build the Neural Network
- Automatic Differentiation with `torch.autograd`
- Optimizing Model Parameters
- Save and Load the Model

Introduction to PyTorch on YouTube [-]

Get Started Ecosystem Mobile Blog **Tutorials** Docs ▾ Resources ▾ GitHub

Tutorials > Learn the Basics

Run in Microsoft Learn Run in Google Colab Download Notebook View on GitHub

Learn the Basics || Quickstart || Tensors || Datasets & DataLoaders || Transforms || Build Model || Autograd || Optimization || Save & Load Model

LEARN THE BASICS

You can play with example codes to learn

Authors: Suraj Subramanian, Seth Juarez, Cassie Breviu, Dmitry Soshnikov, Ari Bornstein

Most machine learning workflows involve working with data, creating models, optimizing model parameters, and saving the trained models. This tutorial introduces you to a complete ML workflow implemented in PyTorch, with links to learn more about each of these concepts.

We'll use the FashionMNIST dataset to train a neural network that predicts if an input image belongs to one of the following classes: T-shirt/top, Trouser, Pullover, Dress, Coat, Sandal, Shirt, Sneaker, Bag, or Ankle boot.

This tutorial assumes a basic familiarity with Python and Deep Learning concepts.



PyTorch Tensor

Tensors are a specialized **data structure** that are very similar to arrays and matrices. In PyTorch, we use tensors to encode the **inputs** and **outputs** of a model, as well as the model's **parameters**.

Tensors can be initialized in various ways.

- Directly from data:

```
data = [[1, 2],[3, 4]]  
x_data = torch.tensor(data)
```

- From a NumPy array:

```
np_array = np.array(data)  
x_np = torch.from_numpy(np_array)
```

- From another tensor:

```
x_ones = torch.ones_like(x_data) # retains the properties of x_data
```

- With random or constant values:

```
shape = (2,3,)  
rand_tensor = torch.rand(shape)  
ones_tensor = torch.ones(shape)  
zeros_tensor = torch.zeros(shape)
```

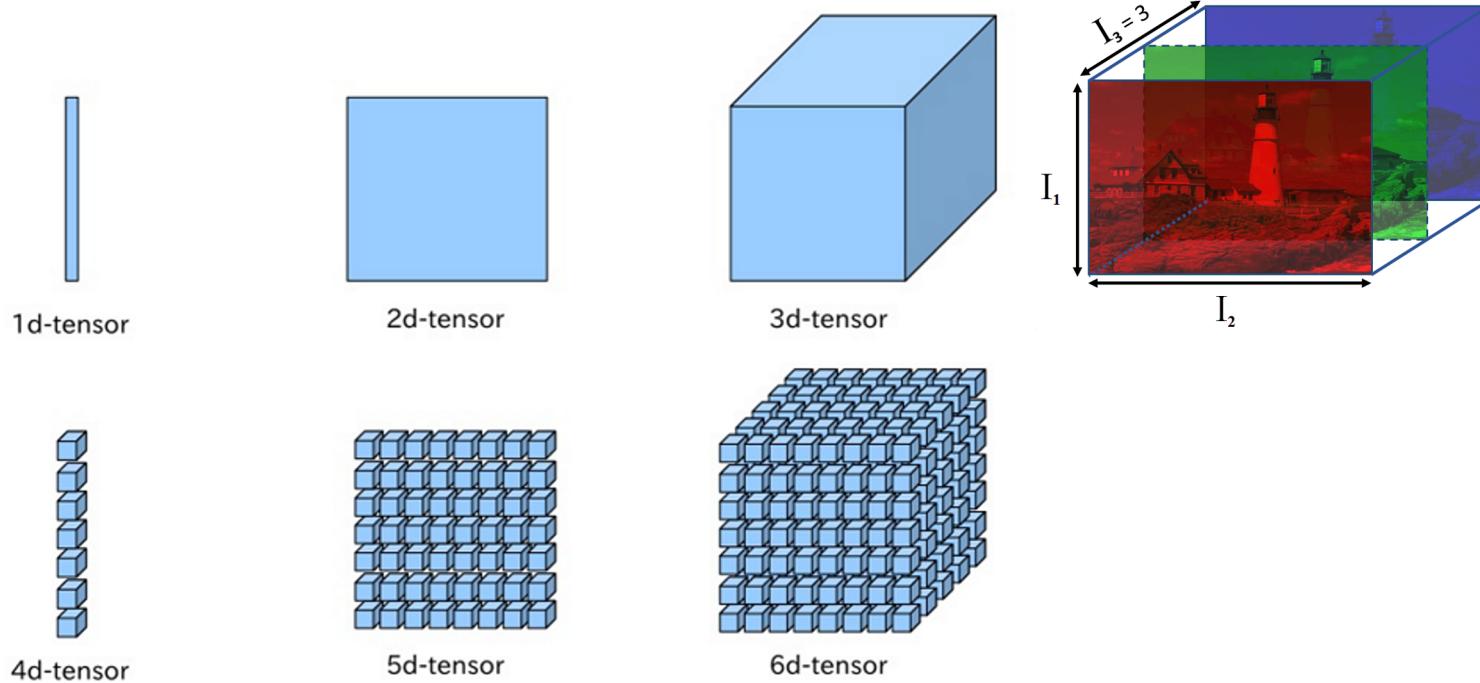


Aside: Tensors

Tensors are multi-dimensional arrays with a uniform data type (called dtype)

- 1d-tensor is a vector, 2d-tensor is a matrix, 3d-tensor is a cube
- 2D color image = 3D tensor

Tensors are similar to numpy arrays, with the addition being that Tensors can also be used on a [GPU](#) to accelerate computing.





Neural network layer in NumPy

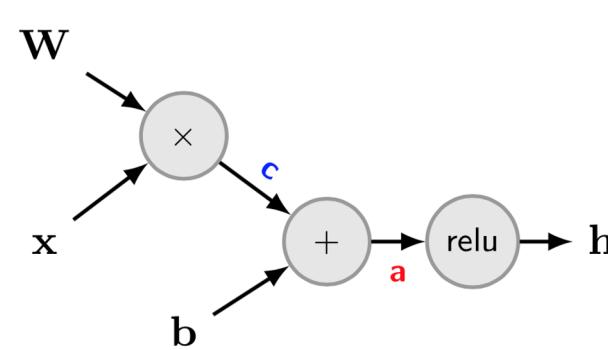
```
import numpy as np

num_in = 10
num_hidden = 50

# init numpy arrays
x = np.random.randn(num_in) # (10,)
W = np.random.randn(num_hidden, num_in) # (50, 10)
b = np.random.randn(num_hidden) # (50,)

# forward pass
a = np.matmul(W, x) + b # get affine output
h = np.maximum(a, np.zeros_like(a))
```

- This is what we're used to.
- It instantiates the following computational graph.





Neural network layer in NumPy

```
import numpy as np

num_in = 10
num_hidden = 50

# init numpy arrays
x = np.random.randn(num_in) # (10,)
W = np.random.randn(num_hidden, num_in) # (50, 10)
b = np.random.randn(num_hidden) # (50,)

# forward pass
a = np.matmul(W, x) + b # get affine output
h = np.maximum(a, np.zeros_like(a))
```

- We can think of each of these functions as implementing some forward pass that operates on inputs.
- e.g., the function `np.matmul(x,y)` has a defined forward pass that takes as input two matrices, `x` and `y`, and multiplies them together.
- But what about a backwards pass?



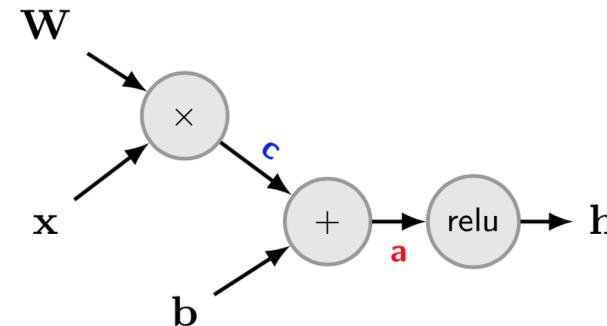
Neural network layer in NumPy

```
import numpy as np

num_in = 10
num_hidden = 50

# init numpy arrays
x = np.random.randn(num_in) # (10,)
W = np.random.randn(num_hidden, num_in) # (50, 10)
b = np.random.randn(num_hidden) # (50,)

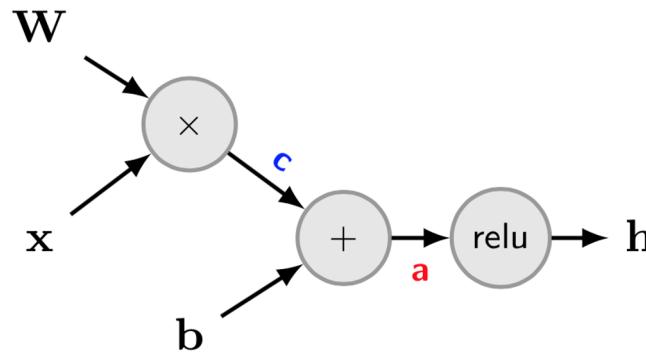
# forward pass
a = np.matmul(W, x) + b # get affine output
h = np.maximum(a, np.zeros_like(a))
```



- As you're used to in the HW, in addition to writing the forward pass, you now have to write a backward pass



Neural network layer in NumPy



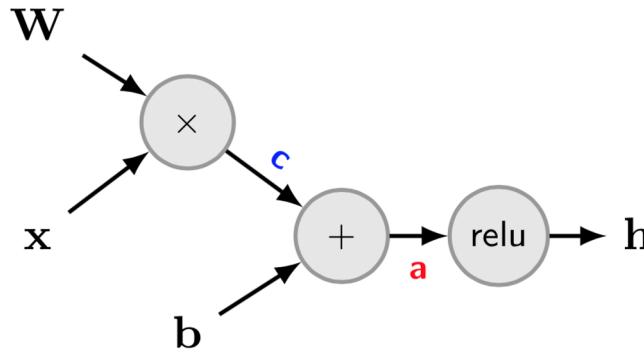
$$\begin{aligned}\frac{\partial \ell}{\partial a} &= \mathbb{I}(a > 0) \odot \frac{\partial \ell}{\partial h} \\ \frac{\partial \ell}{\partial c} &= \frac{\partial a}{\partial c} \frac{\partial \ell}{\partial a} = \frac{\partial \ell}{\partial a} \\ \frac{\partial \ell}{\partial x} &= \frac{\partial c}{\partial x} \frac{\partial \ell}{\partial c} = W^T \frac{\partial \ell}{\partial c} = W^T \frac{\partial \ell}{\partial a} \\ \frac{\partial \ell}{\partial W} &= \frac{\partial c}{\partial W} \frac{\partial \ell}{\partial c} = \frac{\partial \ell}{\partial c} x^T = \frac{\partial \ell}{\partial a} x^T\end{aligned}$$

- Manually implement the backward pass.

```
# backward pass
dL = 1
da = np.expand_dims((a > 0) * dL, axis = 1) # (50,1)
dW = np.matmul(da, np.expand_dims(x, axis = 0)) # (50,1)*(1,10) = (50,10)
dx = np.matmul(W.T, da) # (10,50)*(50,1) = (10,1)
```



The computational graph



$$\begin{aligned}\frac{\partial \ell}{\partial a} &= \mathbb{I}(a > 0) \odot \frac{\partial \ell}{\partial h} \\ \frac{\partial \ell}{\partial c} &= \frac{\partial a}{\partial c} \frac{\partial \ell}{\partial a} = \frac{\partial \ell}{\partial a} \\ \frac{\partial \ell}{\partial x} &= \frac{\partial c}{\partial x} \frac{\partial \ell}{\partial c} = W^T \frac{\partial \ell}{\partial c} = W^T \frac{\partial \ell}{\partial a} \\ \frac{\partial \ell}{\partial W} &= \frac{\partial c}{\partial W} \frac{\partial \ell}{\partial c} = \frac{\partial \ell}{\partial c} x^T = \frac{\partial \ell}{\partial a} x^T\end{aligned}$$

- Couldn't a software **create the computational graph**, store the local gradients, and thus calculate all the gradients with backpropagation for us?
- i.e., in addition to having a “forward” function, it has a “backward” function that accepts the upstream gradient, and multiplies it by the local gradient.
- This is fairly mechanical, and so we should be able to **automate** it.
- This is where deep learning libraries come in.



Implementation in PyTorch

```
import numpy as np
```

```
num_in = 10
num_hidden = 50

# init numpy arrays
x = np.random.randn(num_in)
W = np.random.randn(num_hidden, num_in)
b = np.random.randn(num_hidden)

# forward pass
a = np.matmul(W, x) + b # get affine output
h = np.maximum(a, np.zeros_like(a))
```



```
import torch
```

```
num_in = 10
num_hidden = 50

# init tensors
x = torch.randn(num_in)
W = torch.randn(num_hidden, num_in)
b = torch.randn(num_hidden)

# forward pass
a = torch.matmul(W, x) + b # get affine output
h = torch.max(a, torch.zeros(num_hidden)) # relu
```

- In numpy, `x`, `W`, `b` are all numpy arrays.
- In PyTorch, these are called **PyTorch Tensors**, and they are declared via e.g., `torch.randn`, or `torch.zeros`, etc.
- A PyTorch Tensor is basically the same as a numpy array, but it can run on either CPU or GPU (casted to a `cuda` datatype).



Implementation in PyTorch

Further, it's easy to go between numpy arrays and PyTorch tensors.

```
import torch
import numpy as np

data = [[1, 2],[3, 4]]
np_array = np.array(data)

# numpy array to torch tensor
x_tensor = torch.from_numpy(np_array)

# torch tensor to numpy array
x_array = x_tensor.numpy()

print('x_tensor is:', x_tensor)
print('x_array is', x_array)

x_tensor is: tensor([[1, 2],
                     [3, 4]])
x_array is [[1 2]
             [3 4]]
```

- Note: The torch Tensor and numpy array will share their **underlying memory** locations, and changing one will change the other!



Implementation in PyTorch

```
import torch

num_in = 10
num_hidden = 50

# init tensors
x = torch.randn(num_in)
W = torch.randn(num_hidden, num_in)
b = torch.randn(num_hidden)

# forward pass
a = torch.matmul(W, x) + b # get affine output
h = torch.max(a, torch.zeros(num_hidden)) # relu
```

- Most of the operations that we have in numpy have PyTorch equivalents; usually look up in documentation.

A screenshot of a search results page. On the left, there is a search bar containing the query "torch.randn()". Below the search bar are three navigation links: "PyTorch Recipes [+]", "Introduction to PyTorch [-]", and "Learn the Basics". On the right, a large red-bordered box contains the search results. The results start with the text "Search finished, found 19 page(s) matching the search query." followed by a bulleted list: "• Introduction to PyTorch". Below the list, there is a snippet of text: "...and Long will be the most common. You can create a tensor with random data and the supplied".

- But now PyTorch also, for each of these functions, implements its **backward pass** so it can **automatically** do backprop and calculate gradients for you.



Implementation in PyTorch

```
import torch

num_in = 10
num_hidden = 50

# init tensors
x = torch.randn(num_in)
W = torch.randn(num_hidden, num_in)
b = torch.randn(num_hidden)

# forward pass
a = torch.matmul(W, x) + b # get affine output
h = torch.max(a, torch.zeros(num_hidden)) # relu
```



```
import torch

num_in = 10
num_hidden = 50

# init tensors
x = torch.randn(num_in, requires_grad=True)
W = torch.randn(num_hidden, num_in, requires_grad=True)
b = torch.randn(num_hidden, requires_grad=True)

# forward pass
a = torch.matmul(W, x) + b # get affine output
h = torch.max(a, torch.zeros(num_hidden)) # relu
```

- There is a core torch package called **autograd** for automatic differentiation.
- **requires_grad=True** signals that every operation on the tensor should be tracked.



Implementation in PyTorch

```
num_in = 10
num_hidden = 50

# init tensors with grad
x = torch.randn(num_in, requires_grad=True)
W = torch.randn(num_hidden, num_in, requires_grad=True)
b = torch.randn(num_hidden, requires_grad=True)

# forward pass
a = torch.matmul(W, x) + b # get affine output
h = torch.max(a, torch.zeros(num_hidden)) # relu

print(x.data)
print(x.grad)

tensor([ 0.0156, -1.2956,  1.4077, -0.2591, -0.6188, -1.0828,  0.7685, -0.6054,
        1.2413,  1.0060])
None
```

- `x.data` contains the values of the PyTorch Tensor `x`.
- `x.grad` is another PyTorch **Tensor** containing the gradients of `x`.
- Therefore, `x.grad.data` contains the values of the gradients of `x`.



Implementation in PyTorch

```
num_in = 10
num_hidden = 50

# init tensors with grad
x = torch.randn(num_in, requires_grad=True)
W = torch.randn(num_hidden, num_in, requires_grad=True)
b = torch.randn(num_hidden, requires_grad=True)

# forward pass
a = torch.matmul(W, x) + b # get affine output
h = torch.max(a, torch.zeros(num_hidden)) # relu

# calculate loss
loss = (torch.sum(h) - 1).pow(2)

# do backward
loss.backward()

print("x size: {}, x.grad size: {}".format(list(x.data.shape), list(x.grad.data.shape)))
print("W size: {}, W.grad size: {}".format(list(W.data.shape), list(W.grad.data.shape)))
print("b size: {}, b.grad size: {}".format(list(b.data.shape), list(b.grad.data.shape)))
```

x size: [10], x.grad size: [10]
W size: [50, 10], W.grad size: [50, 10]
b size: [50], b.grad size: [50]

- “`.backward()`” takes the loss function and calculates the gradient for every variable that had a “`requires_grad=True`”.
- This is one line code, i.e., the backward process required no cognitive effort!



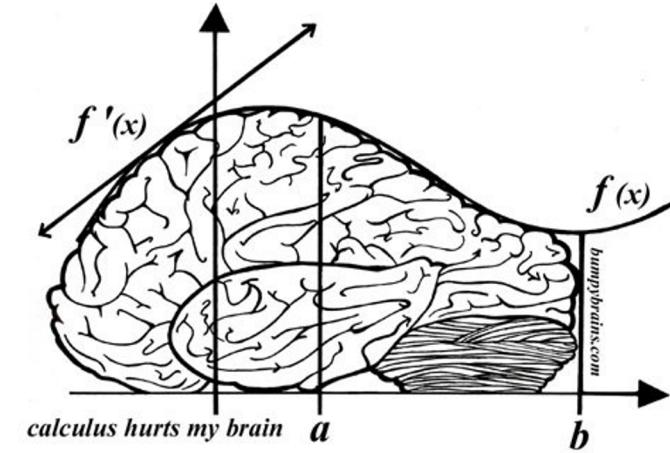
Implementation in PyTorch

Why did we make you implement backpropagation?

It lets you know the call to “`.backward()`” is not magic:

- making a computational graph
- knowing the local gradients at each node of the computational graph
- multiplying them by the upstream gradient

 PyTorch
Autograd

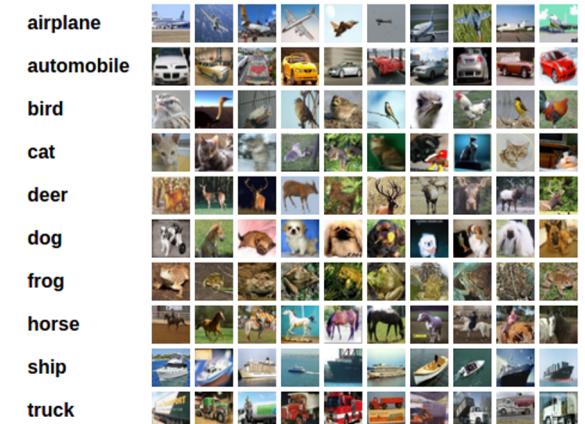
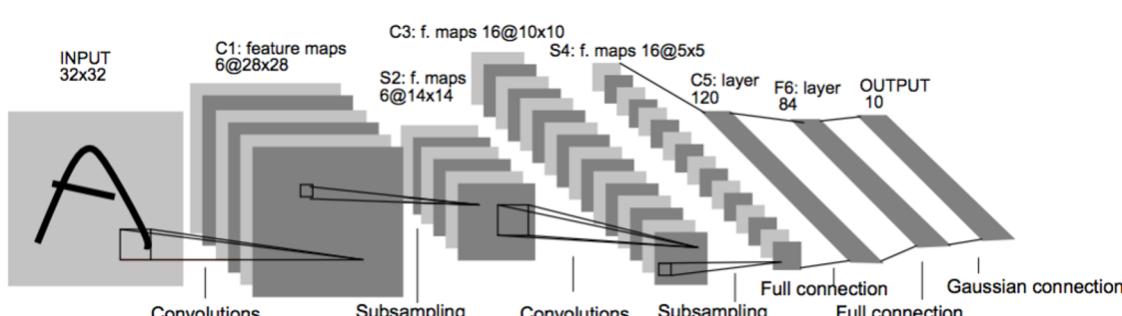




Example: an image classifier

Now with the basic ideas of PyTorch down, we'll go into detail about how to implement a simple neural network.

- Network: the LeNet-5 neural network architecture
- Dataset: CIFAR10, 10 classes, RGB images of size 3x32x32





Example: an image classifier

To learn an image classifier, we will do the following steps in order:

1. Load and normalize the CIFAR10 training and test datasets using [torchvision](#)
2. Define a Convolutional Neural Network
3. Define a loss function
4. Train the network on the training data
5. Test the network on the test data



Step 1: Load and normalize the dataset

Using `torchvision`, it's extremely easy to load CIFAR10.

```
import torch
import torchvision
import torchvision.transforms as transforms
```

```
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

Define the image transformation

```
batch_size = 4
```

```
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                          shuffle=True, num_workers=2)
```

Define the data loader

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                         shuffle=False, num_workers=2)
```

```
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

More details: <https://pytorch.org/vision/stable/transforms.html>



Step 2: Define a convolutional neural network

Modify the LeNet-5 neural network architecture to take 3-channel images

```
import torch.nn as nn  
import torch.nn.functional as F
```

Import torch functions and modules

```
class Net(nn.Module):
```

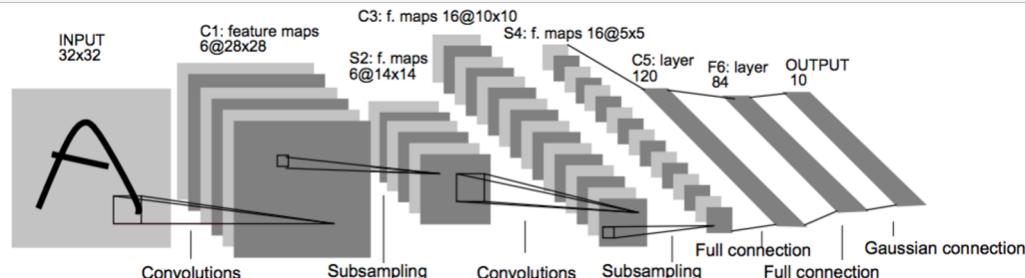
```
    def __init__(self):  
        super().__init__()  
        self.conv1 = nn.Conv2d(3, 6, 5)  
        self.pool = nn.MaxPool2d(2, 2)  
        self.conv2 = nn.Conv2d(6, 16, 5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)
```

Initialize the variables and network modules

```
    def forward(self, x):  
        x = self.pool(F.relu(self.conv1(x)))  
        x = self.pool(F.relu(self.conv2(x)))  
        x = torch.flatten(x, 1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return x
```

Define the forward pass

```
net = Net()
```





Step 2: Define a convolutional neural network

PyTorch incorporates **modules**, which are roughly equivalent to neural network layers.

- A **module** accepts Tensors as inputs and outputs.
- A **module** will also store internal Tensors with learnable weights.
- Examples of **modules** are **Conv2d**, **BatchNorm2d**, i.e., they implement the layers of a neural network.

The package **torch.nn** contains several modules that implement neural network layers, which you can compose together to make networks.

You can essentially think of the nn package as a high-level wrapper that makes it simple to write neural nets.



Step 2: Define a convolutional neural network

All documentation is available online, e.g.

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1,
                     padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros',
                     device=None, dtype=None) [SOURCE]
```

Applies a 2D convolution over an input signal composed of several input planes.

Parameters

- **in_channels** (*int*) – Number of channels in the input image
- **out_channels** (*int*) – Number of channels produced by the convolution
- **kernel_size** (*int* or *tuple*) – Size of the convolving kernel
- **stride** (*int* or *tuple*, *optional*) – Stride of the convolution. Default: 1
- **padding** (*int*, *tuple* or *str*, *optional*) – Padding added to all four sides of the input.
Default: 0
- **padding_mode** (*string*, *optional*) – 'zeros', 'reflect', 'replicate' or
'circular'. Default: 'zeros'
- **dilation** (*int* or *tuple*, *optional*) – Spacing between kernel elements. Default: 1
- **groups** (*int*, *optional*) – Number of blocked connections from input channels to
output channels. Default: 1
- **bias** (*bool*, *optional*) – If `True`, adds a learnable bias to the output. Default: `True`

More details: <https://pytorch.org/docs/stable/nn.html>



Step 2: Define a convolutional neural network

In practice, we can inspect the networks with python functions.

```
print(net)

Net(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
```

The learnable parameters of a net are returned by `net.parameters()`.

```
params = list(net.parameters())
print(len(params))
print(params[0].size()) # conv1's .weight

10
torch.Size([6, 3, 5, 5])
```



Step 3: Define a loss function and optimizer

There are several different loss functions under the `nn` package:

- `nn.MSELoss`
- `nn.CrossEntropyLoss`

`torch.optim` is a package implementing various optimization algorithms:

- `optim.SGD`
- `optim.Adam`
- `optim.RMSprop`

Let's use a classification cross-entropy loss and SGD with momentum.

```
import torch.optim as optim

# define the cross-entropy loss
criterion = nn.CrossEntropyLoss()

# create the optimizer: SGD with momentum
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```



Step 4: Train the network

- We usually utilize scripts similar as follows for training and updating the net.
- We simply have to loop over our data iterator, and feed the inputs to the network and optimize.

```
for epoch in range(10): # loop over the dataset multiple times  
  
    running_loss = 0.0  
    for i, data in enumerate(trainloader, 0):  
        # get the inputs; data is a list of [inputs, labels]  
        inputs, labels = data  
  
        # forward pass  
        outputs = net(inputs)  
        loss = criterion(outputs, labels)  
  
        # backward + optimize  
        loss.backward() # backward tp get gradient values  
        optimizer.step() # does the update  
  
        # zero the parameter gradients  
        optimizer.zero_grad()  
  
        # accumulate loss  
        running_loss += loss.item()
```

Update net, get grad values, and zero grad



Step 4: Train the network

Question: where should I place `optimizer.zero_grad()`?

```
# zero the gradients  
optimizer.zero_grad()
```

```
# forward pass  
outputs = net(inputs)  
loss = criterion(outputs, labels)  
  
# backward + optimize  
loss.backward() # backward  
optimizer.step() # update
```

```
# forward pass  
outputs = net(inputs)  
loss = criterion(outputs, labels)
```

```
# zero the gradients  
optimizer.zero_grad()  
  
# backward + optimize  
loss.backward() # backward  
optimizer.step() # update
```

```
# forward pass  
outputs = net(inputs)  
loss = criterion(outputs, labels)
```

```
# backward + optimize  
loss.backward() # backward  
optimizer.step() # update
```

```
# zero the gradients  
optimizer.zero_grad()
```

All approaches are valid for the standard use case, i.e. if you do not want to accumulate gradients for multiple iterations.

You can thus call `optimizer.zero_grad()` everywhere in the loop but not between the `loss.backward()` and `optimizer.step()` operation.



Step 5: Test the network on the test data

At test time, we set `torch.no_grad()` to disables gradient calculation.

Disabling gradient calculation is useful for inference, when you are sure that you will not call `Tensor.backward()`. It will reduce memory consumption for computations.

```
correct = 0
total = 0

# since we're not training, we don't need to calculate the gradients for our outputs
with torch.no_grad():

    for data in testloader:
        images, labels = data

        # calculate outputs by running images through the network
        outputs = net(images)

        # the class with the highest energy is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f'Accuracy of the network on the 10000 test images: {100 * correct // total}%')
```



Printing the summary of a model

Use the [torchsummary](#) package to print the summary of a model in PyTorch.

```
# !pip3 install torch-summary
from torchsummary import summary
summary(net, (3, 32, 32))
```

```
=====
Layer (type:depth-idx)           Output Shape      Param #
=====
|-Conv2d: 1-1                   [-1, 6, 28, 28]    456
|-MaxPool2d: 1-2                [-1, 6, 14, 14]   --
|-Conv2d: 1-3                   [-1, 16, 10, 10]  2,416
|-MaxPool2d: 1-4                [-1, 16, 5, 5]    --
|-Linear: 1-5                  [-1, 120]          48,120
|-Linear: 1-6                  [-1, 84]           10,164
|-Linear: 1-7                  [-1, 10]            850
=====
Total params: 62,006
Trainable params: 62,006
Non-trainable params: 0
Total mult-adds (M): 0.65
=====
Input size (MB): 0.01
Forward/backward pass size (MB): 0.05
Params size (MB): 0.24
Estimated Total Size (MB): 0.30
=====
```

More details: <https://github.com/sksq96/pytorch-summary>



Saving and loading models

Save/Load `state_dict` (recommended)

- When saving a model for inference, it is only necessary to save the trained model's learned parameters.
- Remember that you must call `model.eval()` to set dropout and batch normalization layers to evaluation mode before running inference.

```
PATH = './cifar_net.pth'

# Save:
torch.save(net.state_dict(), PATH) # save state_dict

# Load:
model = Net()
model.load_state_dict(torch.load(PATH)) # load state_dict
model.eval() # sets model in evaluation (inference) mode
```

More details: https://pytorch.org/tutorials/beginner/saving_loading_models.html



Training on GPU

Define our device as the first visible cuda device if we have CUDA available:

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

# Assuming that we are on a CUDA machine, this should print a CUDA device:

print(device)

cuda:0
```

Transfer the data Tensor and the neural net onto the GPU:

```
net.to(device) # send the net model to GPU

for epoch in range(2): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):

        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data[0].to(device), data[1].to(device) # send data to GPU

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step() # Does the update

        # print statistics
        running_loss += loss.item()
```



Training on GPU

If you have CUDA available, you alternatively use `.cuda()`:

```
net.cuda() # send the net model to GPU

for epoch in range(2): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):

        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data[0].cuda(), data[1].cuda() # send data to GPU

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step() # Does the update

        # print statistics
        running_loss += loss.item()
```



How to adjust learning rate

`torch.optim.lr_scheduler` provides several methods to adjust the learning rate based on the number of epochs:

- `lr_scheduler.ExponentialLR`: decays learning rate by gamma every epoch
- `lr_scheduler.ReduceLROnPlateau`: reduce learning rate when a metric has stopped improving

Learning rate scheduling should be applied after optimizer's update.

```
optimizer = SGD(model, 0.1)
scheduler = ExponentialLR(optimizer, gamma=0.9) # Define the scheduler

for epoch in range(20):
    for input, target in dataset:
        optimizer.zero_grad()
        output = model(input)
        loss = loss_fn(output, target)
        loss.backward()
        optimizer.step()

    scheduler.step() # Learning rate scheduling
```



Getting pre-trained models.

Import `torchvision` to get pre-defined models.

```
import torchvision

# load pre-defined models
alexnet = torchvision.models.alexnet(pretrained=False)

# modify the pre-defined model
my_resnet101 = torchvision.models._dict_['resnet101'](pretrained=False, num_classes=10) # 10 classes
```

Import `torchvision` to get models with pre-trained parameters.

```
import torchvision

# load pre-pretrained models
alexnet = torchvision.models.alexnet(pretrained=True)
vgg16 = torchvision.models.vgg16(pretrained=True)
resnet152 = torchvision.models.resnet152(pretrained=True)

# modify the pre-defined model
model = torchvision.models.resnet101(pretrained=True)
# get features after the last AvgPool
modules = list(model.children())[:-1] # output of pool5: [2048, 1, 1]
model = torch.nn.Sequential(*modules)
```



Build your module

You can combine multiple layers into one module:

```
class FCNet(nn.Module):
    def __init__(self, dims):
        super(FCNet, self).__init__()

        layers = []
        for i in range(len(dims)-2):
            in_dim = dims[i]
            out_dim = dims[i+1]
            layers.append(nn.Linear(in_dim, out_dim))
            layers.append(nn.ReLU())
            layers.append(nn.Dropout(p=0.5))
        layers.append(nn.Linear(dims[-2], dims[-1]))
        layers.append(nn.ReLU())
        layers.append(nn.Dropout(p=0.5))

        self.main = nn.Sequential(*layers)

    def forward(self, x):
        return self.main(x)

fcs = FCNet([16 * 5 * 5, 120, 84, 10])
print(fcs)
```

FCNet(
 (main): Sequential(
 (0): Linear(in_features=400, out_features=120, bias=True)
 (1): ReLU()
 (2): Dropout(p=0.5, inplace=False)
 (3): Linear(in_features=120, out_features=84, bias=True)
 (4): ReLU()
 (5): Dropout(p=0.5, inplace=False)
 (6): Linear(in_features=84, out_features=10, bias=True)
 (7): ReLU()
 (8): Dropout(p=0.5, inplace=False)
)
)



Defining new autograd function

What if the function you want to implement is not in PyTorch?

You can write it yourself: defining the forward and backward pass, and then use it.

Example:

- A third order polynomial, trained to predict $y=g(x)$.

$$P_3(x) = \frac{1}{2}(5x^3 - 3x)$$

- We implement our own custom autograd function

$$P'_3(x) = \frac{3}{2}(5x^2 - 1)$$

More details: https://pytorch.org/tutorials/beginner/examples_autograd/two_layer_net_custom_function.html



Defining new autograd function

This implementation computes the forward pass using operations on PyTorch Tensors, and uses PyTorch autograd to compute gradients.

```
class LegendrePolynomial3(torch.autograd.Function):
    """
    We can implement our own custom autograd Functions by subclassing
    torch.autograd.Function and implementing the forward and backward passes
    which operate on Tensors.
    """

    @staticmethod
    def forward(ctx, input):
        """
        In the forward pass we receive a Tensor containing the input and return
        a Tensor containing the output. ctx is a context object that can be used
        to stash information for backward computation. You can cache arbitrary
        objects for use in the backward pass using the ctx.save_for_backward method.
        """
        ctx.save_for_backward(input)
        return 0.5 * (5 * input ** 3 - 3 * input) →  $P_3(x) = \frac{1}{2}(5x^3 - 3x)$ 

    @staticmethod
    def backward(ctx, grad_output):
        """
        In the backward pass we receive a Tensor containing the gradient of the loss
        with respect to the output, and we need to compute the gradient of the loss
        with respect to the input.
        """
        input, = ctx.saved_tensors
        return grad_output * 1.5 * (5 * input ** 2 - 1) →  $P'_3(x) = \frac{3}{2}(5x^2 - 1)$ 
```

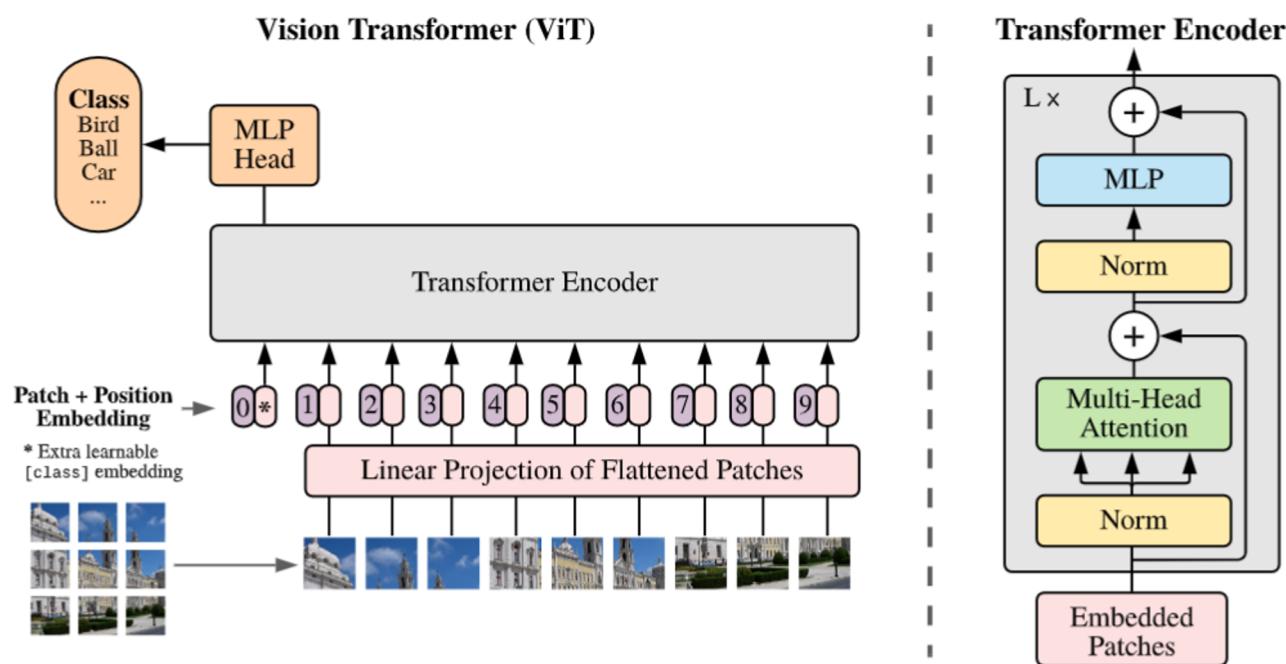


PyTorch Package: Hugging Face



Pretrained models are the state-of-the-art deep learning models.

Hugging Face provides thousands of pretrained models to perform tasks on different modalities such as text, vision, and audio.



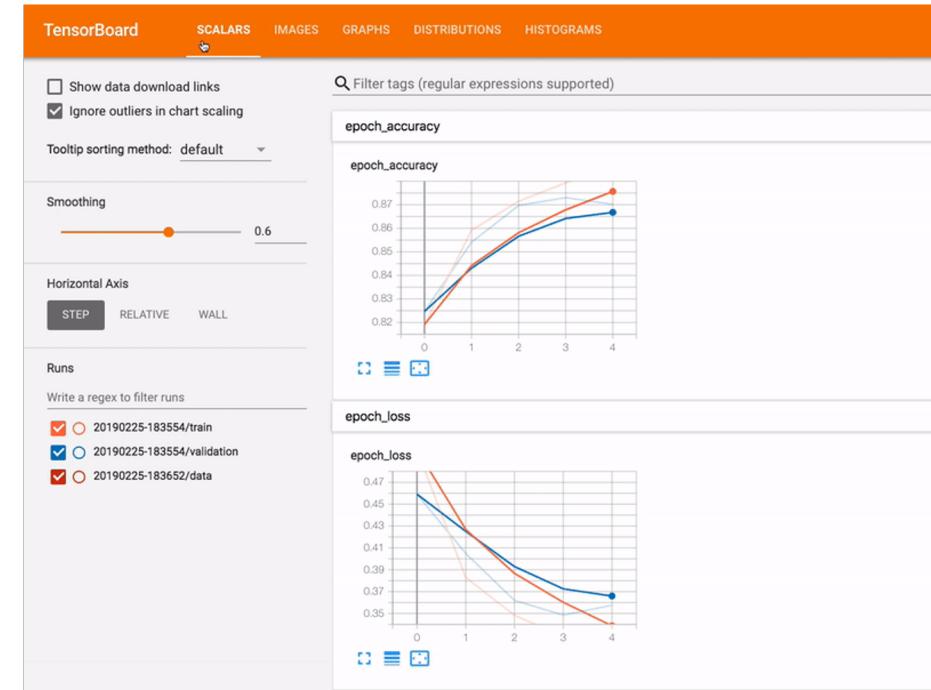
More details: <https://huggingface.co/docs/transformers>



TensorBoard

Visualization toolkit

- Metrics
- Model graphs
- Histograms of weights, biases
- Projecting embeddings
- Displaying images, text
- And much more



More details: https://pytorch.org/tutorials/recipes/recipes/tensorboard_with_pytorch.html



GPU Tool

Monitor the GPU status: `watch nvidia-smi`

Every 2.0s: nvidia-smi											
Sun Feb 20 11:31:33 2022											
GPU			Name		Persistence-M		Bus-Id	Disp.A		Volatile Uncorr. ECC	
Fan	Temp	Perf	Pwr	Usage/Cap				Memory-Usage	GPU-Util	Compute M.	MIG M.
34%	59C	P0	58W / 250W		Off	00000000:01:00.0	On			N/A	
								702MiB / 12192MiB	1%	Default	
										N/A	
0	NVIDIA	TITAN X ...			Off	00000000:01:00.0	On			N/A	
1	NVIDIA	GeForce ...			Off	00000000:02:00.0	Off			N/A	
								10MiB / 24268MiB	0%	Default	
										N/A	
Processes:											
GPU	GI	CI	PID		Type	Process name	GPU Memory		Usage		
ID	ID										
0	N/A	N/A	1689		G	...zoom-client/168/zoom/zoom	34MiB				
0	N/A	N/A	6696		G	/usr/lib/xorg/Xorg	35MiB				



Take-away points

- PyTorch is one of the best deep learning frameworks.
- Pytorch official tutorials
 - <https://pytorch.org/tutorials/>
 - Feel free to check out the tutorial for more details!
- Useful packages
 - CV: torchvision, OpenMMLab
 - NLP: torchtext, Hugging Face, fairseq, AllenNLP
 - Audio: torchaudio
 - Tools: TensorBoard, nvidia-smi, htop
- Some potentially helpful resources:
 - UCB CS285 PyTorch Tutorial: <https://www.youtube.com/watch?v=kPa6hU9prg4>
 - PyTorch for Deep Learning: <https://www.youtube.com/watch?v=GIsq-ZUy0MY>
 - Deep Learning With PyTorch: <https://www.youtube.com/watch?v=c36IUUr864M>
 - <https://d2l.ai/index.html>
 - <https://github.com/bharathgs/Awesome-pytorch-list>



Thanks!
